



Optimierung, komplexe Abfragen und
sharding

Änderung Datenmodell

- Collection **stations** (update on copy)
- Beispiel:

```
{  
  „stationID“ : „23“, // Identifier (string)  
  „creationTS“ : 1314277800000, // Versionierung  
  ... // Metadaten  
}
```

- Abfrage:

```
db.stations.find({„stationID“ : 23, creationTS : { $lt :  
  1314277700000}}).sort({„creationTS“ : -1}).limit(1);
```

Optimierung – document size

- Collection **measureings**
- Beispiel:

```
{
  „timestamp“ : 1314277800000,
  „value“ : 1337,
  „stationID“ : „23“,
  „partID“ : „wr“,
  „serialNo“ : 42,
  „datatype“ : „gain“,
  [„opt1“ : „string“,]
  [„opt2“ : 0]
}
```

```
> db.measureings.stats()
{
  "ns" : "BIS_mongo_eval.measureings",
  "count" : 6000000,
  "size" : 951562096, // coll_size
  "avgObjSize" : 158.59368266666667,
  "storageSize" : 1078239232, // pre_alloc
  "numExtents" : 23,
  "nindexes" : 4,
  "lastExtentSize" : 186253312,
  "paddingFactor" : 1,
  "flags" : 1,
  "totalIndexSize" : 624744512,
  "indexSizes" : {
    "_id_" : 194678736,
    "timestamp_1" : 150953488,
    "datatype_1" : 128158800,
    "value_1" : 150953488
  },
  "ok" : 1
}
```

Optimierung – document size

- Collection **measureings**
- Beispiel:

```
{  
  „a“ : 1314277800000,  
  „b“ : 1337,  
  „c“ : „23“,  
  „d“ : „wr“,  
  „e“ : 42,  
  „f“ : „gain“,  
  [„g“ : „string“,]  
  [„h“ : 0]  
}
```

```
> db.measureings.stats()  
{  
  "ns" : "BIS_mongo_eval.measureings",  
  "count" : 6000000,  
  "size" : 647737272, // coll_size  
  "avgObjSize" : 107.956212,  
  "storageSize" : 891985920, // pre_alloc  
  "numExtents" : 22,  
  "nindexes" : 4,  
  "lastExtentSize" : 155209728,  
  "paddingFactor" : 1,  
  "flags" : 1,  
  "totalIndexSize" : 624744512,  
  "indexSizes" : {  
    "_id_" : 194678736,  
    "a_1" : 150953488,  
    "f_1" : 128158800,  
    "b_1" : 150953488  
  },  
  "ok" : 1  
}
```

Optimierung – document size

- Collection
measurements_<stationID>
- Beispiel:

```
{  
  „a“ : 1314277800000,  
  „b“ : 1337,  
  „d“ : „wr“,  
  „e“ : 42,  
  „f“ : „gain“,  
  [„g“ : „string“,]  
  [„h“ : 0]  
}
```

```
> db.measurements.stats()  
{  
  "ns" : "BIS_mongo_eval.measurements",  
  "count" : 6000000,  
  "size" : 509099848,  
  "avgObjSize" : 84.84997466666667,  
  "storageSize" : 607436800,  
  "numExtents" : 20,  
  "nindexes" : 4,  
  "lastExtentSize" : 107782144,  
  "paddingFactor" : 1,  
  "flags" : 1,  
  "totalIndexSize" : 624744512,  
  "indexSizes" : {  
    "_id_" : 194678736,  
    "a_1" : 150953488,  
    "f_1" : 128158800,  
    "b_1" : 150953488  
  },  
  "ok" : 1  
}
```

Optimierung – document size

- Collection `measurements_<stationID>`
- Beispiel:

```
{  
  „a“ : 1314277800000,  
  „b“ : 1337,  
  „e“ : 42,  
  „f“ : „gain“,  
  [„g“ : „string“,]  
  [„h“ : 0]  
}
```

```
> db.measurements.stats()  
{  
  "ns" : "BIS_mongo_eval.measurements",  
  "count" : 6000000,  
  "size" : 455737340,  
  "avgObjSize" : 75.95622333333333,  
  "storageSize" : 651083776,  
  "numExtents" : 24,  
  "nindexes" : 4,  
  "lastExtentSize" : 111796224,  
  "paddingFactor" : 1,  
  "flags" : 1,  
  "totalIndexSize" : 624744512,  
  "indexSizes" : {  
    "_id_" : 194678736,  
    "a_1" : 150953488,  
    "f_1" : 128158800,  
    "b_1" : 150953488  
  },  
  "ok" : 1  
}
```

Komplexe Abfragen 1

- Wieviele Einträge hat Zeitreihe XY insgesamt/im Zeitintervall [von,bis]?

```
db.measurements.find(  
{  
  "datatype" : <type>,  
  "stationID" : <stationID>,  
  "serialNo" : <serialNo>,  
  "timestamp" :  
  {  
    $gt : <from>,  
    $lt : <to>  
  }  
}).count();
```

```
>db.measurements.find({ "datatype" : "gain", "stationID" :  
  "wendlinghausen2", "serialNo" : 1, "timestamp" : { $gt :  
    1269953100000, $lt : 1269970200000 } }).count();
```

Komplexe Abfragen 2

- Wie ist der Wert der Zeitreihe XY zum Zeitpunkt Z?

```
db.measurements.find(  
{  
  "datatype" : <type>,  
  "stationID" : <stationID>,  
  "serialNo" : <serialNo>,  
  "timestamp" : <t>  
},  
{  
  <projektion> : 1  
});
```

```
> db.measurements.find({ "datatype" : "gain", "stationID" :  
  "wendlinghausen2", "serialNo" : 1, "timestamp" :  
  1269953100000 }, { "value" : 1 } );
```


Komplexe Abfragen 3

- Wie sind die Werte der Zeitreihe XY im Zeitintervall [von,bis]?

```
db.measurements.find(
{
  "datatype" : <type>,
  "stationID" : <stationID>,
  "serialNo" : <serialNo>,
  "timestamp" :
  {
    $gt : <from>,
    $lt : <to>
  }
},
{
  <projektion> : 1
});

>db.measurements.find({ "datatype" : "gain", "stationID" : "wendlinghausen2",
  "serialNo" : 1, "timestamp" : { $gt : 1269953100000, $lt : 1269970200000 }}, {
  "value" : 1 } );
```

Komplexe Abfragen 4

- Wie ist der Zeitpunkt des ältesten/neuesten Eintrags in Zeitreihe XY?

```
db.measurements.find(  
  {  
    "datatype" : <type>,  
    "stationID" : <stationID>,  
    "serialNo" : <serialNo>  
  },  
  {  
    <projektion> : 1  
  }).sort(  
  {  
    <sort_attr> : [1|-1] // should be indexed  
  }).limit(1);
```

```
>db.measurements.find({ "datatype" : "gain", "stationID" :  
  "wendlinghausen2", "serialNo" : 1}, {"timestamp" :  
  1}).sort({"timestamp" : -1}).limit(1); // max
```

Komplexe Abfragen 5 / 1

- Wie ist der maximale/minimale/durchschnittliche Wert der Zeitreihe XY im Zeitintervall [von,bis]?

```
map = function() {  
  emit(this.stationID, // group by  
  {  
    total:this.value,  
    count:1,  
    avg:0,  
    min:this.value,  
    max:this.value  
  });  
}
```

Komplexe Abfragen 5 / 2

```
reduce = function(key, values) {  
  var r = {total:0, count:0, avg:0, min:0, max:0};  
  if(values.length > 0) {  
    r.min = values[0].min;  
    r.max = values[0].max;  
  }  
  values.forEach(function(v) {  
    r.total += v.total;  
    r.count += v.count;  
  
    if(v.min < r.min) {  
      r.min = v.min;  
    }  
    if(v.max > r.max) {  
      r.max = v.max;  
    }  
  });  
  return r;  
}
```

Achtung: for all k,vals : reduce(k, [reduce(k,vals)]) == reduce(k,vals)

Komplexe Abfragen 5 / 3

```
finalize = function(k, r) {  
  if(r.count > 0)  
    r.avg = r.total / r.count;  
  return r;  
}
```

```
db.runCommand(  
{  
  mapreduce : "measurements",  
  map : map,  
  reduce : reduce,  
  out : { inline : 1 },  
  query : { "datatype" : "gain", "stationID" : "wendlinghausen2",  
    "serialNo" : 1, "timestamp" : { $gt : 1269953100000, $lt :  
    1269970200000 } },  
  finalize: finalize  
});
```

Komplexe Abfragen 5 / 4

```
{
  "results" : [
    {
      "_id" : "wendlinghausen2",
      "value" : {
        "total" : 812939,
        "count" : 14,
        "avg" : 58067.07142857143,
        "min" : 46287,
        "max" : 63388
      }
    }
  ],
  "timeMillis" : 34,
  "counts" :
  {
    "input" : 14,
    "emit" : 14,
    "reduce" : 1,
    "output" : 1
  },
  "ok" : 1
}
```

Komplexe Abfragen 6

- Wie ist der Verlauf des Wirkungsgrades für den Wechselrichter XY im Zeitintervall [von,bis]?

```
db.runCommand(  
{  
  mapreduce : "measurements",  
  map : function() {  
    var r = {total_pac:0, total_pdc:0, efficiency:0};  
  
    if(this.datatype == "pac") {  
      r.total_pac = this.value;  
    } else {  
      r.total_pdc = this.value;  
    }  
    emit(this.stationID, r);  
  },  
  reduce : function(key, values) {  
    var r = { total_pac:0, total_pdc:0, efficiency:0 };  
  
    values.forEach(function(v) {  
      r.total_pac += v.total_pac;  
      r.total_pdc += v.total_pdc;  
    });  
  
    return r;  
  },  
  out : { inline : 1 },  
  query : { $or : [ { "datatype" : "pac" }, { "datatype" : "pdc" } ], "stationID" : "wendlinghausen2", "serialNo" :  
    1, "timestamp" : { $gt : 1269953100000, $lt : 1269970200000 } },  
  finalize: function(k, r) {  
    if(r.total_pdc > 0)  
      r.efficiency = r.total_pac / r.total_pdc;  
    return r;  
  }  
});
```

Komplexe Abfragen 7 / 1

- An welchen Tagen hat Zeitreihe XY den Schwellenwert Z über-/unterschritten?

```
db.runCommand(  
{  
  mapreduce : "measureings",  
  map : function() {  
    var r = {date:'', count:1, total:this.value, avg:0, timestamp:this.timestamp}  
    // group documents by day  
    var day = Math.floor(this.timestamp / 1000 / 60 / 60 / 24);  
    emit(day, r);  
  },  
  reduce : function(key, values) {  
    var r = {date:'', count:0, total:0, avg:0, timestamp:0}  
  
    values.forEach(function(v) {  
      r.total += v.total;  
      r.count += v.count;  
      r.timestamp = v.timestamp;  
    });  
    return r;  
  },  
  out : { replace : "temp" },  
  query : { "datatype" : "gain" , "stationID" : "wendlinghausen2", "serialNo" : 1 },  
  finalize: function(k, r) {  
    var date = new Date(r.timestamp);  
    r.date = date.getDate() + "." + (date.getMonth() + 1) + "." + date.getFullYear();  
    if(r.count > 0) {  
      r.avg = r.total / r.count;  
    }  
    return r;  
  }  
});
```


Komplexe Abfragen 7 / 2

```
// query the aggregated data
```

```
db.temp.find({"value.total" : { $lt : 100 } }, { "value.date" : 1 });
```

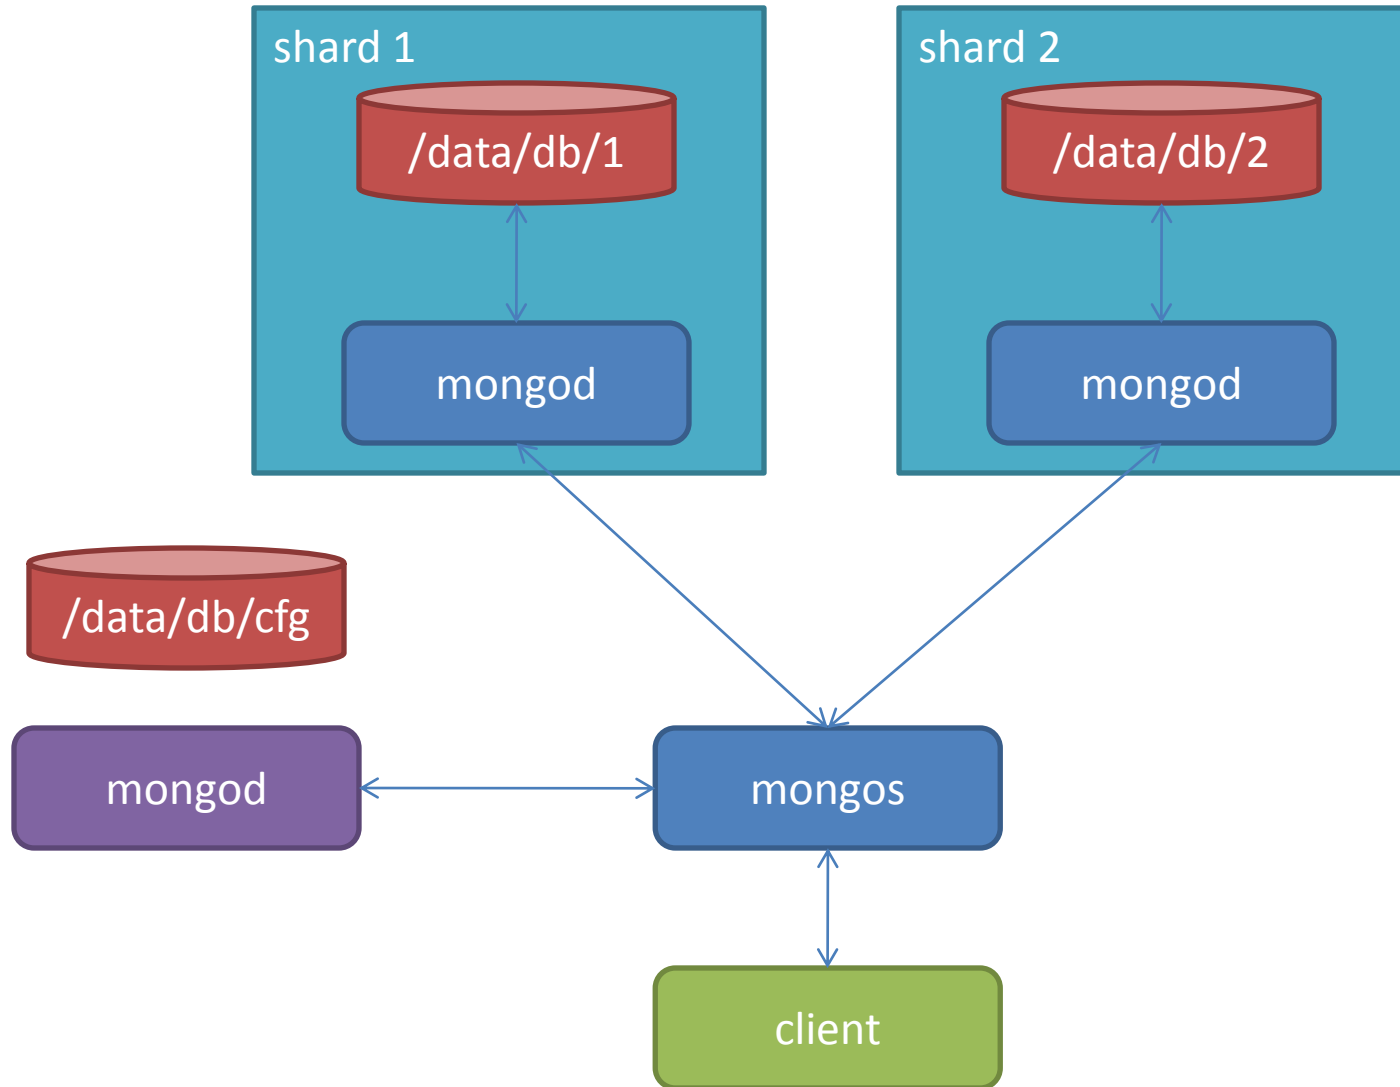
Komplexe Abfragen 8

- Wie groß ist die erzeugte Leistung von Wechselrichter X durchschnittlich pro Temperaturstufe? (d.h. Durchschnitt von PAC je Wert der Temperaturzeitreihe)

Komplexe Abfragen 9

- Für welche Tage im Zeitintervall [von, bis] liegen KEINE Werte für Zeitreihe XY vor?

sharding - Architektur



sharding – server setup

1 the shards

```
> mongod --dbpath=../data/db/1 --port 27017  
> mongod --dbpath=../data/db/2 --port 27018
```

2 the configsrv

```
> mongod --configsvr --dbpath=../data/db/config --port 27019
```

3 mongos

```
> mongos --configdb localhost:27019 --port 27020
```

sharding – sharding setup 1

```
> mongo localhost:27020/admin
```

```
mongos> db.runCommand({addshard : "localhost:27017" });  
{ "shardAdded" : "shard0000", "ok" : 1 }
```

```
mongos> db.runCommand({addshard : "localhost:27018" });  
{ "shardAdded" : "shard0001", "ok" : 1 }
```

```
mongos> db.runCommand( { listshards : 1 } );  
{  
  "shards" : [  
    {  
      "_id" : "shard0000",  
      "host" : "localhost:27017"  
    },  
    {  
      "_id" : "shard0001",  
      "host" : "localhost:27018"  
    }  
  ],  
  "ok" : 1  
}
```

```
mongos> db.runCommand( { enablesharding : "BIS_mongo_eval" } );  
{ "ok" : 1 }
```

```
mongos> db.runCommand({shardcollection : "BIS_mongo_eval.measureings", key : { timestamp : 1 }});  
{ "collectionsharded" : "BIS_mongo_eval.measureings", "ok" : 1 }
```

sharding – sharding setup 2

```
mongos> db.printShardingStatus({verbose : 1})
--- Sharding Status ---
  sharding version: { "_id" : 1, "version" : 3 }
  shards:
    { "_id" : "shard0000", "host" : "localhost:27017" }
    { "_id" : "shard0001", "host" : "localhost:27018" }
  databases:
    { "_id" : "admin", "partitioned" : false, "primary" : "config" }
    { "_id" : "BIS_mongo_eval", "partitioned" : true, "primary" : "shard0000" }
      BIS_mongo_eval.measureings chunks:
        shard0001      10
        shard0000      10
    { "timestamp" : { $minKey : 1 } } --> { "timestamp" : NumberLong("1269953100000") } on : shard0001 { "t" : 8000, "i" : 1 }
    { "timestamp" : NumberLong("1269953100000") } --> { "timestamp" : NumberLong("1303706700000") } on : shard0001 { "t" : 7000, "i" : 2 }
    { "timestamp" : NumberLong("1303706700000") } --> { "timestamp" : NumberLong("1303927200000") } on : shard0001 { "t" : 9000, "i" : 2 }
    { "timestamp" : NumberLong("1303927200000") } --> { "timestamp" : NumberLong("1305022500000") } on : shard0001 { "t" : 9000, "i" : 3 }
    { "timestamp" : NumberLong("1305022500000") } --> { "timestamp" : NumberLong("1305873900000") } on : shard0001 { "t" : 4000, "i" : 0 }
    { "timestamp" : NumberLong("1305873900000") } --> { "timestamp" : NumberLong("1306164600000") } on : shard0001 { "t" : 8000, "i" : 4 }
    { "timestamp" : NumberLong("1306164600000") } --> { "timestamp" : NumberLong("1307038500000") } on : shard0001 { "t" : 8000, "i" : 5 }
    { "timestamp" : NumberLong("1307038500000") } --> { "timestamp" : NumberLong("1307439000000") } on : shard0001 { "t" : 7000, "i" : 4 }
    { "timestamp" : NumberLong("1307439000000") } --> { "timestamp" : NumberLong("1308294000000") } on : shard0001 { "t" : 7000, "i" : 5 }
    { "timestamp" : NumberLong("1308294000000") } --> { "timestamp" : NumberLong("1308827820000") } on : shard0001 { "t" : 9000, "i" : 0 }
    { "timestamp" : NumberLong("1308827820000") } --> { "timestamp" : NumberLong("1309076100000") } on : shard0000 { "t" : 9000, "i" : 1 }
    { "timestamp" : NumberLong("1309076100000") } --> { "timestamp" : NumberLong("1310128260000") } on : shard0000 { "t" : 8000, "i" : 9 }
    { "timestamp" : NumberLong("1310128260000") } --> { "timestamp" : NumberLong("1310752800000") } on : shard0000 { "t" : 6000, "i" : 2 }
    { "timestamp" : NumberLong("1310752800000") } --> { "timestamp" : NumberLong("1311094800000") } on : shard0000 { "t" : 8000, "i" : 2 }
    { "timestamp" : NumberLong("1311094800000") } --> { "timestamp" : NumberLong("1312012800000") } on : shard0000 { "t" : 8000, "i" : 3 }
    { "timestamp" : NumberLong("1312012800000") } --> { "timestamp" : NumberLong("1312211700000") } on : shard0000 { "t" : 9000, "i" : 4 }
    { "timestamp" : NumberLong("1312211700000") } --> { "timestamp" : NumberLong("1313221500000") } on : shard0000 { "t" : 9000, "i" : 5 }
    { "timestamp" : NumberLong("1313221500000") } --> { "timestamp" : NumberLong("1313510400000") } on : shard0000 { "t" : 8000, "i" : 6 }
    { "timestamp" : NumberLong("1313510400000") } --> { "timestamp" : NumberLong("1314469800000") } on : shard0000 { "t" : 8000, "i" : 7 }
    { "timestamp" : NumberLong("1314469800000") } --> { "timestamp" : { $maxKey : 1 } } on : shard0000 { "t" : 1000, "i" : 4 }
```

todo

- Sort Internals (interne Optimierung)
- PAC / PDC für jeden Timestamp (Verlauf)
- „Date“ in Query 7 aus map und reduce rausnehmen
- Komplexe Anfragen 8 und 9
- MapReduce Chunk Definition / Konfiguration
- Sharding Internals näher erläutern
- 1 Milliarde Zeitreihen (duplizieren?, echte Daten?) -> gemeinsames testbed
- Benchmarking (ebenfalls gemeinsames testbed, sonst keine Aussagen möglich)
 - Map Reduce auf sharded DB / single DB
 - JSON Queries auf sharded DB / single DB
- Schriftliche Ausarbeitung zum Vergleich GraphDB / Document Store
- Abschlusspräsentation vorbereiten